


MPC Scaling Efficiency: Per-Eval Cost Amortization



CUDA Graphs: scaling approaches pure kernel time. Eager: scales slower due to launch overhead.