


MPC Rollouts (256 evals): Optimization Effectiveness Comparison



Key insight: KV Cache hurts MPC (many evals amortize prefill). CUDA Graphs shine (short compute). Combined: compute approaches pure kernel time.