


Cache Efficiency: Where Is Data Actually Coming From?



CUDA Graphs improve cache locality → higher hit rates. KV Cache reduces L2 hits (introduces scattered loads).